

Greedy Algorithms

Greedy algorithms is another useful way for solving optimization problems.

Optimization Problems

- For the given input, we are seeking solutions that must satisfy certain conditions.
- These solutions are called feasible solutions. (In general, there are many feasible solutions.)
- We have an optimization measure defined for each feasible solution.
- We are looking for a feasible solution that optimizes (either maximum or minimum) the optimization measure.

Examples

Matrix Chain Product Problem

- A **feasible solution** is any valid parenthesization of an n -term chain.
- The **optimization measure** is the total number of scalar multiplications for the parenthesization.
- Goal: Minimize the the total number of scalar multiplications.

0/1 Knapsack Problem

- A **feasible solution** is any subset of items whose total weight is at most the knapsack capacity K .
- The **optimization measure** is the total item profit of the subset.
- Goal: Maximize the the total profit.

General Description

- Given an optimization problem P , we seek an optimal solution.
- The solution is obtained by a sequence of steps.
- In each step, we select an “item” to be included into the solution.
- At each step, the decision is made based on the selections we have already made so far, that looks the best choice for achieving the optimization goal.
- Once a selection is made, it cannot be undone: The selected item cannot be removed from the solution.

Minimum Spanning Tree (MST) Problem

This is a classical graph problem. We will study graph algorithms in detail later. Here we use MST as an example of Greedy Algorithms.

Definition

A **tree** is a **connected** graph with **no cycles**.

Definition

Let $G = (V, E)$ be a graph. A **spanning tree** of G is a subgraph of G that **contains all vertices of G and is a tree**.

Minimum Spanning Tree (MST) Problem

Input: An connected undirected graph $G = (V, E)$. Each edge $e \in E$ has a **weight** $w(e) \geq 0$.

Find: a spanning tree T of G such that $w(T) = \sum_{e \in T} w(e)$ **is minimum**.

Kruskal's Algorithm


Kruskal's Algorithm

- 1: Sort the edges by non-decreasing weight. Let e_1, e_2, \dots, e_m be the sorted edge list
- 2: $T \leftarrow \emptyset$
- 3: **for** $i = 1$ **to** m **do**
- 4: **if** $T \cup \{e_i\}$ does not contain a cycle **then**
- 5: $T \leftarrow T \cup \{e_i\}$
- 6: **else**
- 7: do nothing
- 8: **end if**
- 9: **end for**
- 10: **output** T

Kruskal's Algorithm

- The algorithm goes through a sequence of steps.
- At each step, we consider the edge e_i , and decide whether add e_i into T .
- Since we are building a spanning tree T , T can not contain any cycle. So if adding e_i into T introduces a cycle in T , we do not add it into T .
- Otherwise, we add e_i into T . We are processing the edges in the order of increasing edge weight. So when e_i is added into T , it looks the best to achieve the goal (minimum total weight).
- Once e_i is added, it is never removed and is included into the final tree T .
- This is a perfect example of greedy algorithms.

An Example



- The number near an edge is its weight. The **blue edges** are in the MST constructed by Kruskal's algorithm.
- The **blue numbers in ()** indicate the order in which the edges are added into MST.

Kruskal's Algorithm

- For a given graph $G = (V, E)$, its MST is not unique. However, the weight of any two MSTs of G must be the same.
- In Kruskal's algorithm, two edges e_i and e_{i+1} may have the same weight. If we process e_{i+1} before e_i , we may get a different MST.
- Runtime of Kruskal's algorithm:
 - Sorting of edge list takes $\Theta(m \log m)$ time.
 - Then we process the edges one by one. So the loop iterates m time.
 - When processing an edge e_i , we check if $T \cup \{e_i\}$ contains a cycle or not. If not, add e_i into T . If yes, do nothing.
 - By using proper data structures, the processing of an edge e_i can be done in $O(\log n)$ time. (The detail was discussed in CSE250).
 - So the loop takes $O(m \log n)$ time.
 - Since G is connected, $m \geq n$. The total runtime is $\Theta(m \log m + m \log n) = \Theta(m \log m)$.

Elements of Greedy Algorithms

- Are we done?
- **No!** A big task is not done yet: How do we know Kruskal's algorithm is correct?
- Namely, how do we know the tree constructed by Kruskal's algorithm is indeed a MST?
- You may have **convinced** yourself that we are using an **obvious** strategy towards the optimization goal.
- In this case, we are lucky: our intuition is correct.
- But in other cases, the strategies that **seem equally obvious** may lead to wrong solutions.
- In general, the correctness of a greedy algorithm **requires proof**.

Correctness Proof of Algorithms

- An algorithm A is **correct**, if it works on **all** inputs.
- If A works on some inputs, but not on some other inputs, then A is **incorrect**.
- To show A is **correct**, you must argue that **for all inputs, A produces intended solution**.
- To show A is **incorrect**, you only need to give a **counter example input I** : You show that, for this particular input I , the output from A is not the intended solution.
- Strictly speaking, all algorithms need correctness proof.
- For DaC, it's often so straightforward that the correctness proof is unnecessary/omitted. (Example: MergeSort)
- For dynamic programming algorithms, the correctness proof is less obvious than the DaC algorithms. But in most time, it is quite easy to **convince** people (**i.e. informal proof**) the algorithm is correct.
- For greedy algorithms, the correctness proof can be very **tricky**.

Elements of Greedy Algorithms

For a greedy strategy to work, it must have the following two properties.

Optimal Substructure Property

An optimal solution of the problem contains within it the optimal solutions of subproblems.

- This is the same property required by the dynamic programming algorithms.

Greedy Choice Property

A global optimal solution can be obtained by making a locally optimal choice that seems the best toward the optimization goal when the choice is made. (Namely: The choice is made based on the choices we have already made, **not** based on the future choices we might make.)

- This property is harder to describe exactly.
- Best way to understand it is by examples.



Optimal Substructure Property for MST

Example

Optimal Substructure Property for MST

- Let $G = (V, E)$ be a connected graph with edge weight.
- Let $e_1 = (x, y)$ be the edge with the smallest weight. (Namely, e_1 is the first edge chosen by Kruskal's algorithm.)
- Let $G' = (V', E')$ be the graph obtained from G by merging x and y :
 - x and y becomes a single new vertex z in G' .
 - Namely $V' = V - \{x, y\} \cup \{z\}$
 - e_1 is deleted from G .
 - Any edge e_i in G that was incident to x or y now is incident to z .
 - The edge weights remain unchanged.

Optimal Substructure Property for MST



Optimal Substructure Property for MST

If T is a MST of G containing e_1 , then $T' = T - \{e_1\}$ is a MST of G' .


Greedy Choice Property for Kruskal's Algorithm

- Let e_1, e_2, \dots, e_m be the edge list in the order of increasing weight. So e_1 is the first edge chosen by Kruskal's algorithm.
- Let T_{opt} be an MST of G . By definition, the total weight of T_{opt} is the minimum.
- We want to show T_{opt} contains e_1 .
- But this is not always possible. Recall that the MST of G is not unique.
- So we will do this: Starting from T_{opt} , we change T_{opt} , without increasing the weight in the process, to another MST T' that contains e_1 .
- If T_{opt} contains e_1 , then we are done (lucky!)

Greedy Choice Property for Kruskal's Algorithm

- Suppose T_{opt} does not contain e_1 .
- Consider the graph $H = T_{opt} \cup \{e_1\}$.
- H contains a cycle C . Let $e_i \neq e_1$ be another edge on C .
- Let $T' = T_{opt} - \{e_i\} \cup \{e_1\}$.
- Then T' is a spanning tree of G .
- Since e_1 is the edge with the smallest weight, $w(e_1) \leq w(e_i)$.
- Hence $w(T') = w(T_{opt}) - w(e_i) + w(e_1) \leq w(T_{opt})$.
- But T_{opt} is a MST!
- So we must have $w(e_i) = w(e_1)$ and $w(T_{opt}) = w(T')$. In other words, both T_{opt} and T' are MSTs of G .
- This is what we want to show: **There is an MST that contains e_1 . So when Kruskal's algorithm includes e_1 into T , we are not making a mistake.**

Greedy Choice Property for Kruskal's Algorithm



Correctness Proof of Kruskal's Algorithm

- The proof is by induction.
- Kruskal's algorithm selects the lightest edge $e_1 = (x, y)$.
- By **Greedy Choice Property**, there exists an optimal MST of G that contains e_1 .
- By induction hypothesis, Kruskal's algorithm construct a MST T' in the graph $G' = ((V - \{x, y\} \cup \{z\}), E')$ which is obtained from G by merging the two end vertices x, y of e_1 .
- By the **Optimal Substructure Property of MST**, $T = T' \cup \{e_1\}$ is a MST of G .
- This T is the tree constructed by Kruskal's algorithm. Hence, Kruskal's algorithm indeed returns a MST.

0/1 Knapsack Problem

We mentioned that **some seemingly intuitive greedy strategies do not really work**. Here is an example.

0/1 Knapsack Problem

Input: n item $_i$ ($1 \leq i \leq n$). Each item $_i$ has an integer weight $w[i] \geq 0$ and a profit $p[i] \geq 0$.

A knapsack with an integer capacity K .

Find: A subset of items so that the total weight of the selected items is at most K , and the total profit is maximized.

There are several **greedy strategies that seem reasonable**. But none of them works.

0/1 Knapsack Problem

Greedy Strategy 1

Since the goal is to maximize the profit without exceeding the capacity, we fill the items in the order of increasing weights. Namely:

- Sort the items by increasing item weight: $w[1] \leq w[2] \leq \dots$.
- Fill the knapsack in the order $\text{item}_1, \text{item}_2, \dots$ until no more items can be put into the knapsack without exceeding the capacity.

Counter Example:

$n = 2, w[1] = 2, w[2] = 4, p[1] = 2, p[2] = 3, K = 4$.

- This strategy puts item_1 into the knapsack with total profit 2.
- The optimal solution: put item_2 into the knapsack with total profit 3.

0/1 Knapsack Problem

- For this greedy strategy, we can still show the **Optimal Substructure Property** holds:
 - if S is an optimal solution, that contains the item_1 , for the original input,
 - then $S - \{\text{item}_1\}$ is an optimal solution for the input consisting of $\text{item}_2, \text{item}_3, \dots, \text{item}_n$ and the knapsack with capacity $K - w[1]$.
- However, we cannot prove the **Greedy Choice Property**: We are not able to show there is an optimal solution that contains the item_1 (the lightest item).
- **Without this property, there is no guarantee this strategy would work.** (As the counter example has shown, it doesn't work.)

0/1 Knapsack Problem

Greedy Strategy 2

Since the goal is to maximize the profit without exceeding the capacity, we fill the items in the order of decreasing profits. Namely:

- Sort the items by decreasing item profit: $p[1] \geq p[2] \geq \dots$.
- Fill the knapsack in the order item₁, item₂, ... until no more items can be put into the knapsack without exceeding the capacity.

Counter Example:

$n = 3, p[1] = 3, p[2] = 2, p[3] = 2, w[1] = 3, w[2] = 2, w[3] = 2, K = 4$.

- This strategy puts item₁ into the knapsack with total profit 3.
- The optimal solution: put item₂ and item₃ into the knapsack with total profit 4.

0/1 Knapsack Problem

Greedy Strategy 3

Since the goal is to maximize the profit without exceeding the capacity, we fill the items in the order of decreasing **unit profit**. Namely:

- Sort the items by decreasing item unit profit: $\frac{p[1]}{w[1]} \geq \frac{p[2]}{w[2]} \geq \frac{p[3]}{w[3]} \dots$
- Fill the knapsack in the order $\text{item}_1, \text{item}_2, \dots$ until no more items can be put into the knapsack without exceeding the capacity.

Counter Example:

$n = 2, w[1] = 2, w[2] = 4, p[1] = 2, p[2] = 3, K = 4.$

- We have: $\frac{p[1]}{w[1]} = \frac{2}{2} = 1 \geq \frac{p[2]}{w[2]} = \frac{3}{4}.$
- This strategy puts item_1 into knapsack with total profit 2.
- The optimal solution: put item_2 into knapsack with total profit 3.

Fractional Knapsack Problem

Fractional Knapsack Problem

Input: n items i ($1 \leq i \leq n$). Each item i has an integer weight $w[i] \geq 0$ and a profit $p[i] \geq 0$.

A knapsack with an integer capacity K .

Find: A subset of items to put into the knapsack. **We can select a fraction of an item.** The goal is the same: the total weight of the selected items is at most K , and the total profit is maximized.

Mathematical description of Fractional Knapsack Problem

Input: $2n + 1$ integers $p[1], p[2], \dots, p[n], w[1], w[2], \dots, w[n], K$

Find: a vector (x_1, x_2, \dots, x_n) such that:

- $0 \leq x_i \leq 1$ for $1 \leq i \leq n$
- $\sum_{i=1}^n x_i \cdot w[i] \leq K$
- $\sum_{i=1}^n x_i \cdot p[i]$ is maximized.

Fractional Knapsack Problem

- Although the **Fractional Knapsack Problem** looks very similar to the **0/1 Knapsack Problem**, it is much much easier.
- The **Greedy Strategy 3** works.

Greedy-Fractional-Knapsack

- 1: Sort the items by decreasing unit profit: $\frac{p[1]}{w[1]} \geq \frac{p[2]}{w[2]} \geq \frac{p[3]}{w[3]} \dots$
- 2: $i = 1$
- 3: **while** $K > 0$ **do**
- 4: **if** $K > w[i]$ **then**
- 5: $x_i = 1$ **and** $K = K - w[i]$
- 6: **else**
- 7: $x_i = K/w[i]$ **and** $K = 0$
- 8: **end if**
- 9: $i = i + 1$
- 10: **end while**

It can be shown the Greedy Choice Property holds in this case.

Activity Selection Problem

Activity Selection Problem

- A set $S = \{1, 2, \dots, n\}$ of **activities**.
- Each activity i has a **starting time** s_i and a **finishing time** f_i ($s_i \leq f_i$).
- Two activities i and j are **compatible** if the interval $[s_i, f_i)$ and $[s_j, f_j)$ do not **overlap**.
- Goal: Select a subset $A \subseteq S$ of **mutually compatible activities** so that $|A|$ is maximized.

Application

- Consider a single CPU computer. It can run only one job at any time.
- Each activity i is a job to be run on the CPU that must start at time s_i and finish at time f_i .
- How to select a maximum subset A of jobs to run on CPU?

Greedy Algorithm for Activity Selection Problem


Greedy Strategy

At any moment t , select the activity i with the smallest finish time f_i .


Greedy-Activity-Selection

- 1: Sort the activities by increasing finish time: $f_1 \leq f_2 \leq \dots \leq f_n$
- 2: $A = \{1\}$ (A is the set of activities to be selected.)
- 3: $j = 1$ (j is the current activity being considered.)
- 4: **for** $i = 2$ **to** n **do**
- 5: **if** $s_i \geq f_j$ **then**
- 6: $A = A \cup \{i\}$
- 7: $j = i$
- 8: **end if**
- 9: **end for**
- 10: return A

Example



Input



After Sorting

Solid lines are selected activities

Dashed lines are not selected

- $[1, 3)$ is the first interval selected. The dashed intervals $[0, 4)$ and $[2, 6)$ are killed because they are not compatible with $[1, 3)$.
- This problem is also called the interval scheduling problem.

Proof of Correctness

- Let $S = \{1, 2, \dots, n\}$ be the set of activities to be selected. Assume $f_1 \leq f_2 \leq \dots \leq f_n$.
- Let O be an optimal solution. Namely O is a subset of mutually compatible activities and $|O|$ is maximum.
- Let X be the output from the Greedy algorithm. We always have $1 \in X$.
- We want to show $|O| = |X|$. We will do this by induction on n .

Greedy Choice Property

- The activity 1 is selected by the greedy algorithm. We need to show there is an optimal solution that contains the activity 1.
- If the optimal solution O contains 1, we are done.
- If not, let k be the first activity in O . Let $O' = O - \{k\} \cup \{1\}$.
- Since $f_1 \leq f_k$, all activities in O' are still mutually compatible.
- Clearly $|O| = |O'|$. So O' is an optimal solution containing 1.


Proof of Correctness

By the **Greedy Choice Property**, we may assume the optimal solution O contains the job 1.

Optimal Substructure Property

- Let $S_1 = \{i \in S \mid s_i \geq f_1\}$. (S_1 is the set of jobs that are compatible with job 1. Or equivalently, the set of jobs that are not killed by job 1.)
- Let $O_1 = O - \{1\}$.
- Claim: O_1 is an optimal solution of the job set S_1 .
 - If this is not true, let O'_1 be an optimal solution set of S_1 . Since O_1 is not optimal, we have $|O'_1| > |O_1|$.
 - Let $O' = O'_1 \cup \{1\}$. Then O' is a set of mutually compatible jobs in S , and $|O'| = |O'_1| + 1 > |O_1| + 1 = |O|$.
 - But O is an optimal solution. This is a contradiction.
- Hence the claim is true.

Proof of Correctness



Proof of Correctness

- Since the **Optimal Substructure** and **Greedy Choice** properties are true, we can prove the correctness of the greedy algorithm by induction.
- Greedy algorithm picks the job 1 in its solution.
- By the **Greedy Choice property**, there is an optimal solution that also contains the job 1. So this selection needs not be reversed.
- The greedy algorithm delete all jobs that are incompatible with job 1. The remaining jobs is the set S_1 in the proof of **Optimal Substructure** property.
- By induction hypothesis, Greedy algorithm will output an optimal solution X_1 for S_1 .
- By the **Optimal Substructure** property, $X = X_1 \cup \{1\}$ is an optimal solution of the original job set S .
- X is the output from Greedy algorithm. So the algorithm is correct.
- Runtime: Clearly $O(n \log n)$ (dominated by sorting).


Scheduling All Intervals

- Schedule all activities using as few resources as possible.
- **Input:**
 - A set $\mathcal{R} = \{I_1, \dots, I_n\}$ of n requests/activities.
 - Each I_i has a start time s_i and finish time f_i . (So each I_i is represented by an interval $[s_i, f_i)$).
- **Output:** A partition of \mathcal{R} into as few subsets as possible, so that the intervals in each subset are mutually compatible. (Namely, they do not overlap.)

Scheduling All Intervals

Application


- Each request I_i is a job to be run on a CPU.
- If two intervals I_p and I_q overlap, they cannot run on the same CPU.
- How to run all jobs using as few CPUs as possible?



Scheduling All Intervals

Another way to look at the problem:

- Color the intervals in \mathcal{R} by different colors.
- The intervals with the same color do not overlap.
- Using as few colors as possible.




This problem is also known as **Interval Graph Coloring Problem**.

Scheduling All Intervals

Graph Coloring

Let $G = (V, E)$ be an undirected graph.

- A **vertex coloring** of G is an assignment of **colors** to the vertices of G so that no two vertices with the same color are adjacent to each other in G .
- Equivalently, a **vertex coloring** of G is a partition of V into vertex subsets so that no two vertices in the same subset are adjacent to each other.



A vertex coloring is also called just **coloring** of G . If G has a coloring with k colors, we say G is **k -colorable**.

Scheduling All Intervals

Graph Coloring Problem

Input: An undirected graph $G = (V, E)$

Output: Find a vertex coloring of G using as few colors as possible.

Chromatic Number

$\chi(G)$ = the smallest k such that G is k -colorable



- $\chi(G) = 1$ iff G has no edges.
- $\chi(G) = 2$ iff G is a bipartite graph with at least 1 edge.
- Graph Coloring is a **very hard** problem.
- The problem can be solved in poly-time **only for special graphs**.

Scheduling All Intervals

Four Color Theorem

Every **planar graph** can be colored using at most 4 colors.

G is a planar graph if it can be drawn on the plane so that no two edges cross.





Both graphs (a) and (b) are planar graphs. The graph (a) has a 3-coloring. The graph (b) **requires** 4 colors, because all 4 vertices are adjacent to each other, and hence each vertex must have a different color.

Scheduling All Intervals

Interval Graph

$G = (V, E)$ is called an **interval graph** if it can be represented as follows:

- Each vertex $p \in V$ represents an interval $[b_p, f_p)$.
- $(p, q) \in E$ if and only if the two intervals I_p and I_q overlap.



Scheduling All Intervals



- It is easy to see that the problem of **scheduling all intervals** is precisely the **graph coloring problem for interval graphs**.
- We discuss a greedy algorithm for solving this problem.
- It is not easy to prove the **greedy choice property** for this greedy strategy.
- We show the correctness of the algorithm by other methods.
- We use **queues** Q_1, Q_2, \dots to **hold** the subsets of intervals. (You can think that each Q_i is a CPU, and if an interval $I_p = [b_p, f_p)$ is put into Q_i , the job p is run on that CPU.)
- Initially all queues are empty.
- When we consider an interval $[b_p, f_p)$ and a queue Q_i , we look at the last interval $[b_t, f_t)$ in Q_i . If $f_t \leq b_p$, we say **Q_i is available for $[b_p, f_p)$** . (Meaning: the CPU Q_i has finished the last job assigned to it. So it is ready to run the job $[b_p, f_p)$.)

Scheduling All Intervals


Greedy-Schedule-All-Intervals

- 1 sort the intervals according to increasing b_p value: $b_1 \leq b_2 \leq \dots \leq b_n$
- 2 $k = 0$ (k will be the number of queues we need.)
- 3 **for** $p = 1$ **to** n **do**:
- 4 look at Q_1, Q_2, \dots, Q_k , put $[b_p, f_p)$ into the first available Q_i .
- 5 if no current queue is available:
 - increase k by 1;
 - open a new empty queue;
 - put $[b_p, f_p)$ into this new queue.
- 6 **output** k and Q_1, \dots, Q_k

Scheduling All Intervals



After Sorting



Q1

Q2

Q3

Scheduling All Intervals

Proof of correctness:

- We only put intervals into available queues. So each queue contains only non-overlapping intervals.
- We need to show the algorithm uses **minimum number of queues**. (Namely, partition intervals into minimum number of subsets.)
 - If the input contains k mutually overlapping intervals, we must use at least k queues. (Because no two such intervals can be placed into the same queue.)
 - When the algorithm opens a new empty queue Q_k for an interval $[b_p, f_p)$, none of the current queues Q_1, \dots, Q_{k-1} is available. This means that the last intervals in Q_1, \dots, Q_{k-1} all overlap with $[b_p, f_p)$. Hence the input contains k mutually overlapping intervals.
 - The algorithm uses k queues. By the observation above, this is the smallest possible.

Scheduling All Intervals

Runtime Analysis:

- Sorting takes $O(n \log n)$ time.
- The loop runs n times.
- The loop body scans Q_1, \dots, Q_k to find the first available queue. So it takes $O(k)$ time.
- Hence, the runtime is $\Theta(nk)$, (where k is the number of queues needed, or equivalently the chromatic number $\chi(G)$ of the input interval graph G .)

In the worst case, k can be $\Theta(n)$. Hence, the worst case runtime is $\Theta(n^2)$.